



HEALTHCARE ANALYSIS: RESPIRATORY ILLNESS

PRESENTED BY THE S3 AMIGOS

JASMINE JIANG | ETTIONE STUCKEY II | JUSTIN QUINN

TABLE OF CONTENTS

01 Overview

03 Setting up S3

05 Lambda to RDS

07 Hosting on EC2

02 Schema

04 Setting up RDS

06 RDS to Analysis

08 Lessons Learned

+

YIKES!

Respiratory illness is an ever-growing problem, threatening the lives of Americans all across the country ...





YAY!

The good news is, The S3 Amigos are on the case! With our skills in data analytics and AWS, we can recognize trends and patterns in the data and provide a basis on which future healthcare strategies will be developed!

TEAM STRATEGY: DIVIDE & CONQUER



OBJECTIVES

- Create a working pipeline from file upload to cloud RDS
- Clean data via deployed Lambda function
- Upload data to RDS
- Connect locally to RDS to generate analytics
- Display analytics results on EC2-hosted website



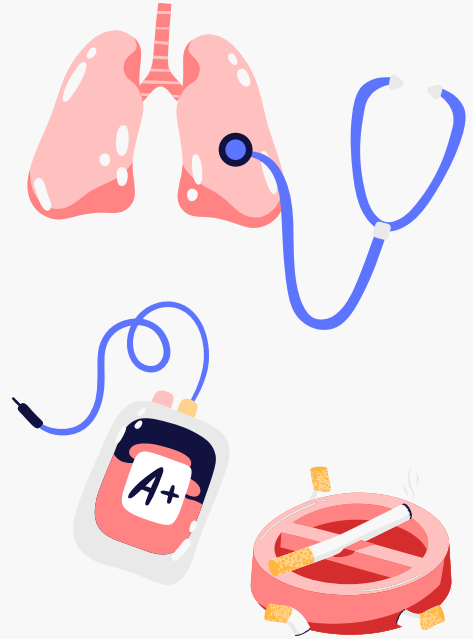
02

SCHEMA

ER DIAGRAM



Outpatient_Respiratory_Illness_Activity	
PK	entry_id INTEGER NOT NULL
	week_ending DATE
	week INTEGER
	season TEXT
	state TEXT
	activity_level TEXT
	activity_level_label TEXT



AWS ARCHITECTURE

Let's take a look at how our AWS services are set up and connected to each other.



03

CHECKPOINT 1: SETTING UP THE S3



CONFIGURATION: S3

Select a region

AWS Region

US East (N. Virginia) us-east-1

Create a name

Bucket name [Info](#)

myNewBucket

Bucket name must be unique within the global namespace and follow the bucket naming rules. [See rules for bucket naming](#)

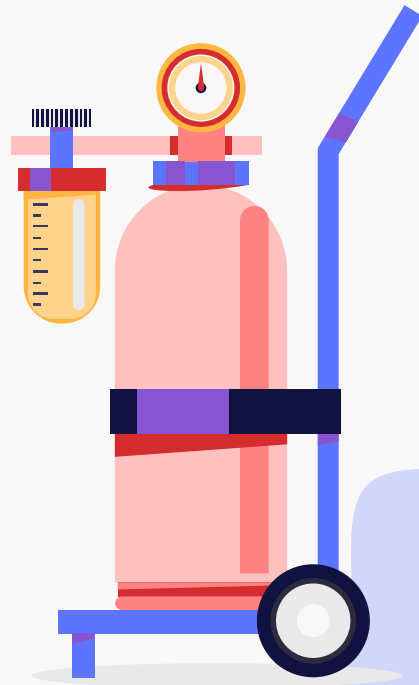
Block public access

☒ **Block all public access**

Turning this setting on is the same as turning on all four settings below. Each of the following settings are independent of one another.

Create bucket

That's it! The S3 has been created.



04

CHECKPOINT 2: SETTING UP THE RDS



CONFIGURATION: RDS

Select an engine & template



Free tier

Use RDS Free Tier to develop new applications, test existing applications, or gain hands-on experience with Amazon RDS. [Info](#)

Set db identifier & credentials

DB instance identifier [Info](#)

Type a name for your DB instance. Region.

The DB instance identifier is case-insensitive characters or hyphens. First character must be a letter.

Master username [Info](#)

Type a login ID for the master user of your DB instance.

1 to 16 alphanumeric characters. The first character must be a letter.

Master password [Info](#)

Constraints: At least 8 printable ASCII characters.

Select the default VPC

Virtual private cloud (VPC) [Info](#)

Choose the VPC. The VPC defines the virtual networking environment for this DB instance.

6 Subnets, 6 Availability Zones

Only VPCs with a corresponding DB subnet group are listed.



WHOA!

Remember to set your RDS public access to “yes”. Failure to do so could cause complications when trying to access the database with other services!

Public access [Info](#)

☒ **Yes**
RDS assigns a public IP address to the database. Amazon EC2 instances and other resources outside of the VPC can connect to your database. Resources inside the VPC can also connect to the database. Choose one or more VPC security groups that specify which resources can connect to the database.

☐ **No**
RDS doesn't assign a public IP address to the database. Only Amazon EC2 instances and other resources inside the VPC can connect to your database. Choose one or more VPC security groups that specify which resources can connect to the database.

Create table



05

CHECKPOINT 3: LAMBDA TO RDS

CONNECTING TO DATABASE

- Created new Lambda function
- Assigned appropriate permissions and security groups
- Used Pymysql library to create connection to RDS database and to clean data
- Attached function to S3 trigger



POPULATING DATABASE

```
bucket = event['Records'][0]['s3']['bucket']['name']
key = urllib.parse.unquote_plus(event['Records'][0]['s3']['object']['key'], encoding='utf-8')
local_file_path = '/tmp/' + os.path.basename(key)
s3.download_file(bucket, key, local_file_path)
logger.info("SUCCESS: File retrieved from S3")

data = []
with open(local_file_path, 'r') as csvfile:
    csv_reader = csv.reader(csvfile)
    for row in csv_reader:
        data.append(row[:4] + [row[4].replace("Level ", '')] + row[5:])
    data = [tuple(row) for row in data]
logger.info("SUCCESS: File read")
data = data[1:]
```



Downloaded .csv file from S3 bucket
into /tmp folder



Cleaned the extracted
data



Executed insert statement with parsed
list of tuples

PERMISSIONS: IAM

Although our resources have now been configured, there are still a few things that need to be done before the team can access them.

- IAM users must be created for all group members.
- IAM users must be added to a group.
- A policy must be created and added to the group.

Users in this group (3)

An IAM user is an entity that you create in AWS to represent the person or application that uses it to interact with AWS.

☐ User name ☐ [Ettione](#)☐ [Jasmine](#)☐ [Justin](#)

```
1 {  
2   "Version": "2012-10-17",  
3   "Statement": [  
4     {  
5       "Sid": "VisualEditor0",  
6       "Effect": "Allow",  
7       "Action": [  
8         "s3:*",  
9         "iam:*",  
10        "ec2:*",  
11        "logs:*",  
12        "q:*",  
13        "ec2-instance-connect:*"  
14      ],  
15      "Resource": [  
16        "*"   
17      ]  
18    },  
19    {  
20      "Sid": "Statement1",  
21      "Effect": "Allow",  
22      "Action": [  
23        "lambda:*",  
24        "rds:*",  
25        "cloudwatch:*"  
26      ],  
27      "Resource": [  
28        "*"   
29      ]  
30    }  
31  ]  
32 }
```



06

CHECKPOINT 4: RDS TO ANALYSIS

Upload to the S3 Bucket

```
#Programmatically upload to S3 bucket
import boto3
import os

aws_access_key = os.environ['ACCESS_KEY']
aws_secret_key = os.environ['SECRET_KEY']

bucket_name = 'pep2-group1-etl-endpoint'

file_key = 'Outpatient_Respiratory_Illness_Activity_Map'

local_file_path = 'Healthcare-Data-Analysis/Outpatient_Respiratory_Illness_Activity_Ma

s3 = boto3.client('s3', aws_access_key_id=aws_access_key, aws_secret_access_key=aws_se
|
with open(local_file_path, 'rb') as data:
    s3.put_object(Body=data, Bucket=bucket_name, Key=file_key)

print(f"File {file_key} uploaded to {bucket_name}")
```



Be sure to keep those keys private!

```
s3 = boto3.client('s3')
user_name = os.environ['USER_NAME']
password = os.environ['PASSWORD']
rds_proxy_host = os.environ['RDS_PROXY_HOST']
db_name = os.environ['DB_NAME']

logger = logging.getLogger()
logger.setLevel(logging.INFO)

try:
    conn = pymysql.connect(host=rds_proxy_host, user=
except pymysql.MySQLError as e:
    logger.error("ERROR: Unexpected error: Could not
    logger.error(e)
    sys.exit(1)
```

Connect to RDS + Upload CSV

Be sure you have the correct RDS Proxy Host and DB Name.

Using Try-Except blocks is good practice!

Pull from RDS + Convert to Dataframe

Double check your queries to ensure everything runs smoothly.

Congrats! You are now ready to do some analysis!

```
import pymysql
import pandas as pd
import os
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

user_name = os.environ['USER_NAME']
password = os.environ['PASSWORD']
rds_proxy_host = os.environ['RDS_PROXY_HOST']
db_name = os.environ['DB_NAME']

try:
    conn = pymysql.connect(host=rds_proxy_host,
                           user=user_name,
                           password=password,
                           database=db_name)
except pymysql.MySQLError as e:
    print(e)

query = 'select * from respirator'
try:
    df = pd.read_sql_query(query, conn)
    df.set_index('entry_id', inplace=True)
except pd.DatabaseError as e:
    print(e)
finally:
    conn.close()
```



07

EXTENSION: HOSTING ON EC2

Creating the EC2

Step 1: Enter EC2 Dashboard



Step 2: Select 'Instances'

▼ Instances

Instances

Step 3: Launch an Instance

Launch instances

Step 4: Select OS

Amazon
Linux

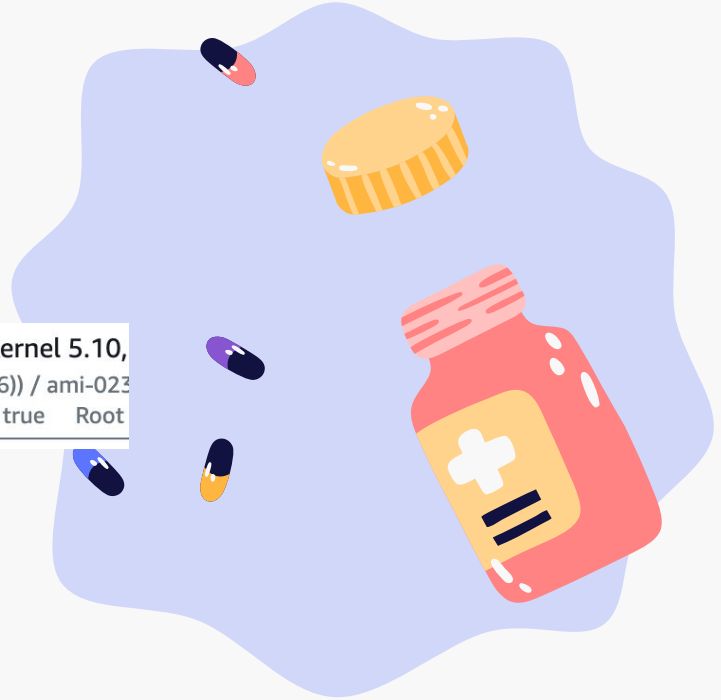


Step 5: AMI

Amazon Linux 2 AMI (HVM) - Kernel 5.10,
ami-0cf10cdf9fcd62d37 (64-bit (x86)) / ami-023
Virtualization: hvm ENA enabled: true Root

Step 6: Architecture

64-bit (x86)



Creating the EC2+ WebHosting

Step 7: Instance Type

Instance type

t2.micro

Family: t2 1 vCPU 1

Step 8: Key Pair

[Create new key pair](#)

Create key pair

Step 9: Launch

Launch instance

Connect

Root User

```
$ sudo su -
```

Install

```
]$ yum install -y httpd
```

Host

```
systemctl enable httpd
```

08

LESSONS LEARNED



CHALLENGE #1

The Problem:

Connection failed during testing despite configured environment variables

Tried multiple potential endpoints for host variable

Finally pinpointed and solved issue by configuring security groups for Lambda function

The Lesson:

Don't get fixated on what you may *think* is the problem, and consider potential outside factors as well

CHALLENGE #2

The Problem:

File from S3 bucket could not be downloaded via Lambda

Created a new Lambda function from scratch, which *could* download and process the file, and compared them

Only difference was that the original function was in a VPC

I unhooked the function from the VPC and set the host to the public RDS address

The Lesson:

Go with what works, even if you don't entirely understand the fine mechanics

Use process of elimination and comparison with working code to pinpoint issues

CHALLENGE #3

The Problem:

Was unable to connect to AWS RDS from local machines.

Reconfigured security groups, proxies, and VPCs to no avail.

Finally, went through our RDS configurations and noticed that public access was set to "No".

After changing this setting to "yes" and waiting a couple minutes for the modification to take effect, we were finally able to connect to the database.

The Lesson:

Sometimes the *simplest* solution is the best solution. Be sure to carefully look over the configurations for your service before it is deployed!

CHALLENGE #4

The Problem:

Was unable to remote connect into the EC2 from a local machine

Even though the key file and EC2 address were correct, my access was denied.

The error message stated my key was not secure enough

This require that I change the permissions on the key file directly through the command line

The Lesson:

Read! Very often the system is telling you exactly what you are doing wrong. Understanding error messages can help point you in the right direction to find a solution.

DEMO TIME!

Jira Board

Website

CREDITS: This presentation template was created by **Slidesgo**, and includes icons by **Flaticon**, infographics & images by **Freepik**

